

From ASTRA 1.0 to ASTRA 2.0

A complete comparative report: what changed, what it buys, what it costs, and what is still unmeasured

Comparison report

18 August 2026

Abstract

ASTRA 1.0 executes one *atomic cycle*: a single decidable proposition, proposed by an ensemble, translated into an executable validator by a different model, reviewed by a third, and run by a deterministic oracle. ASTRA 2.0 keeps that worker unchanged in spirit and adds a layer above it—an append-only evidence ledger, structured proposal portfolios, hard admissibility gates, progressive widening, and resume—so that a research *programme*, not a single question, can be driven to a decision and audited afterwards.

This report states the differences precisely, separates what is already established from what is not, and carries the head-to-head measurement that until 17 August did not exist. On the same 43-record suite, 1.0 accepts 41.7% of false claims and 2.0 accepts none; the difference in *overall* accuracy is not statistically significant, but the false-acceptance difference is, and becomes decisive once the false-claim sample is enlarged—0 of 22 for 2.0 against 11 of 22 for 1.0, exact $p \approx 0.001$. What changed is the kind of error, not the rate of correctness: 1.0 fails by asserting falsehoods and 2.0 fails by running out of time.

1 Executive summary

ASTRA is a research engine that forces every claim through an executable falsification step. **1.0** runs one such step—the *atomic cycle*: an ensemble proposes a decidable proposition, one model authors a validator, a second reviews it, a deterministic oracle runs it. **2.0** keeps that worker and adds a layer above it so that a research *programme*, not a single question, can be driven to a decision and audited afterwards: an append-only evidence ledger, structured proposal portfolios, hard admissibility gates, progressive widening, resume, and—the epistemic core—five status axes that are never collapsed into one verdict.

	ASTRA 1.0	ASTRA 2.0
Unit of work	one atomic cycle	a campaign of cycles over a ledger
Verdict shape	a single status	five separate axes (operation, claim, evidence, branch, goal)
False claims accepted	11 of 22 (50%)	0 of 22 (0%), exact $p \approx 0.001$
Overall accuracy	lower, but not significantly	higher, but not significantly
Dominant error	asserting a falsehood	running out of time
Latency per case	p50 187 s / p95 676 s	p50 327 s / p95 946 s (1.75×)
Record	a transcript that scrolls past	auditable evidence with hashes
Negative results	repeated	cumulative (a closed branch)
Self-inflicted failures	cannot exhibit them	exhibits them, and can catch them

The honest bottom line, stated once and not softened: **2.0 is demonstrably more careful, not demonstrably more accurate.** Its measured, pre-registered advantage is the elimination of *false acceptances*—confident wrong answers—which is the error a downstream reader discovers

only after acting on it. Its cost is latency and the operational complexity of a layer that acts on its own signals. Everything below documents both sides.

2 Scope: this is an architecture for commercial models

ASTRA is built to do research with the models that are *commercially available*—subscription-tier CLIs from more than one provider, on a single workstation, under ordinary quota. It is not built on, and makes no claim about, the private or unreleased higher-capability systems the provider companies hold internally.

That constraint is the premise of the design, not a limitation apologised for. Everything distinctive about the architecture follows from it:

- **Independence is bought by heterogeneity, not by scale.** Three providers with different training and different blind spots give a cross-check that a single stronger model cannot give itself.
- **Refutation is mechanical because the models cannot be trusted to self-assess.** The deterministic auditor, the executable validator and the oracle exist precisely because no available model reliably reports when its own argument fails.
- **Budget discipline is structural.** Phase ceilings, downstream reserves and a machine-wide single-cycle lock exist because quota is scarce and shared, not elastic.

The consequence is worth stating plainly: if a substantially stronger model becomes available, ASTRA does not become unnecessary—it becomes cheaper to run, because fewer repair rounds are needed. The architecture is about *where evidence comes from*, and that question does not dissolve with model capability.

3 The two lines side by side

3.1 What 1.0 is

One bounded investigation, end to end: parallel proposals from two providers, cross-critique, synthesis into a consensus conjecture, validator authoring by a third provider, deterministic preflight, independent review, oracle execution, evidence audit, and a navigator that proposes the next direction. Five status axes are already kept separate. It is a strong single-question engine.

What it does not have is memory across questions. Each cycle begins from a prompt and ends with a verdict and a suggested direction; the minority proposal that lost the synthesis, the branch that was refuted, and the reason either happened are not represented anywhere a later cycle can consult.

3.2 What 2.0 adds

4 Measured differences

4.1 Code

A drift check compares every shared module three ways—against the common base commit, the current 2.0 tree, and the current production tree.

Classification	Files	Meaning
IN SYNC	60	identical behaviour on both lines
BOTH CHANGED	5	diverged on both sides, needs a decision at merge
2.0 ONLY	15	development evolution not yet in production
PRODUCTION ONLY	0	nothing in production is missing from 2.0




Figure 1: 2.0 does not replace the worker; it closes a loop around it. The upward arrow is the part 1.0 has no place to put.

The zero matters: 2.0 is a strict superset of production behaviour plus its own evolution, so a comparison cannot be explained by 2.0 having lost something.

	ASTRA 1.0	ASTRA 2.0
campaign modules (<code>core/campaign_*.py</code>)	0	9
test files	25	54
collected tests	152	469

4.2 Epistemic controls added

Each of these was added because a live measurement exposed the failure it prevents; none is speculative.

Control	Failure it prevents
verdict re-anchoring	a corrected version of a false claim reported as validating the original
decisive vs auxiliary legs	a weak auxiliary check deciding a PASS, or an auxiliary failure printing a scientific FAIL
<code>not_code_at_all</code>	prose handed back to be patched as if it were a damaged script
symbolic cost control	a pointwise claim attacked with full symbolic machinery until memory dies
conjecture fencing	text produced by one agent obeyed as instructions by another
plan-artifact sanitising	a read-only sandbox's scaffolding leaking into the consensus
thinking budget cap	an output ceiling consumed entirely by internal reasoning
downstream budget reserve	the repair step, which runs last, starved by a proportional rule
machine-wide cycle lock	two checkouts deliberating against one model account

4.3 Ledger guarantees added

These have no counterpart in 1.0 because 1.0 has nothing to guarantee: it does not persist a research record.

- append-only events with a canonical-JSON SHA-256 chain;
- atomic checkpoints, and a truncated tail detected and archived rather than silently repaired;
- content-addressed frozen resources, so editing an input is detectable;
- a replayed cycle refused rather than recorded as a second observation;
- a result bound to the request's objective, so a cycle that answered a different question cannot become evidence;
- human interventions recorded as first-class decisions with an actor, rather than silently changing the record.

5 What is not measured yet

5.1 The head-to-head measurement

Both lines were run on the same suite, tier and oracle: 43 records, 23 paired scientific cases, 14 audit, 6 execution.

Standard tier	ASTRA 1.0	ASTRA 2.0
false acceptance	0.4167	0.0
false rejection	0.0	0.0
strict accuracy	0.5217	0.6957
95% CI	[0.330, 0.708]	[0.491, 0.844]
operational failure rate	0.2609	0.2609
audit: critical defect recall	1.0	1.0
audit: false alarm, sound code	0.3333	0.0
execution: verdict accuracy	1.0	1.0
latency p50 / p95	187 / 676 s	327 / 946 s

The result is the kind of error, not the rate of correctness. Over the same 23 cases:

ASTRA 1.0	11 wrong: 5 validated , 5 API_ERROR, 1 CODE_ERROR
ASTRA 2.0	7 wrong: 4 TIMEOUT, 2 API_ERROR, 1 INCONCLUSIVE

The five VALIDATED entries are *false* claims accepted as valid, and 2.0 refuted all five. They are the signature of the re-anchoring defect: production identifies the falsehood correctly, forms a corrected conjecture, validates the correction, and reports the status of *its* conjecture rather than the claim it was asked about. Every step is right and the answer is to a different question.

So 1.0 fails by asserting falsehoods; 2.0 fails by running out of time. A timeout is a budget problem. A false acceptance is a wrong answer delivered with confidence, and it is only discovered by re-reading the paper after it ships.

What the statistics support, and what they do not. With 23 paired cases, the difference in *overall accuracy is not significant*: McNemar’s exact test over the discordant pairs (5 favouring 1.0, 9 favouring 2.0) gives $p = 0.42$, and the confidence intervals overlap widely. We do not claim 2.0 is more accurate.

The false-acceptance difference is 5 of 12 false claims against 0 of 12, with all five discordant pairs in the same direction: exact $p = 0.0625$. Suggestive, narrowly above the conventional threshold, and—worth stating—the smallest p this design can produce with twelve false claims and five discordant pairs.

Powering the false-acceptance test. Because that p was floor-limited by sample size rather than by effect, the suite was extended with ten more false claims and both lines were re-run on them under matching conditions (standard tier, local oracle, `-case-root` so 1.0 sees the same cases). The claims are representative false versions of true theorems across logic, GR, quantum, fluids and ODE, and were *not* chosen to trigger the defect. 1.0 accepted **6 of the 10** as valid; 2.0 accepted none, refuting eight and timing out on two. Six discordant pairs, all favouring 2.0, give exact $p = 0.03125$ on the new cases alone—below the conventional threshold; pooled with the original five, eleven discordant pairs all in one direction give exact $p \approx 0.001$. The pre-registered false-acceptance advantage is now significant.

What that number does and does not license. The new claims lean toward the *near-miss* type—false but close to a true statement—which is both the realistic case and precisely the regime where the re-anchoring defect bites, so the raw strict-accuracy gap they produce (1 of 10 for 1.0 against 8 of 10 for 2.0) is inflated by that selection and is *not* a general accuracy claim. What survives cleanly is the specific one: on false claims, 1.0 accepts about half (11 of 22 across both runs) and 2.0 accepts none (0 of 22). 2.0’s two failures here are timeouts and three of 1.0’s are operational errors—neither a false acceptance—and both are excluded from the discordant count, which is over acceptance alone.

Two caveats belong with the table. The audit case file differs between the checkouts by a fixture correction, so that row is not strictly paired; the 23 scientific cases are (verified identical modulo line endings). And 1.0 was measured on today’s production checkout, which already carries the two fixes ported on 13 August, so this is not the July production line.

5.2 What this measurement did not settle

One run per line, so the operational failure rates agreeing exactly (0.2609 on both) is evidence of comparable conditions rather than proof of them, and nothing here has been repeated. 2.0 also costs $1.75\times$ the wall time per case, which is the price of more review and more repair and shows up directly in its errors: four of its seven failures are timeouts.

5.3 The benchmark cannot see the main contribution

This is the honest limit of any 1.0/2.0 comparison. The quality benchmark scores single cycles. The campaign layer is not exercised by it at all, and 1.0 has no campaign layer to compare against, so on that axis the difference is not a number but a presence.

What the campaign layer buys has to be argued from what it makes possible, and so far the evidence is qualitative but concrete: in live use it recorded refutations with exact witnesses against a persistent claim set, promoted a branch on negative evidence under a deterministic rule, refused to widen on an episode that never ran, retracted a contaminated record without rewriting history, and stopped a validator whose conclusions came from undefined variables. None of those events has a place to happen in 1.0.

5.4 And a first comparative signal that is not favourable

Within 2.0, a preregistered ablation of the reflective loop against a matched linear control did not meet its own interpretability criterion—operational failure stayed far above the 0.2 threshold fixed in advance—and the raw counts leaned against the reflective arm. That result is reported in the architecture report rather than hidden here. It concerns the *policy* inside 2.0, not the 1.0/2.0 difference, but anyone reading this document should know it exists.

6 Why the extra structure is worth it

Three arguments, in decreasing order of how well they are evidenced.

It makes negative results usable. Most research time is spent pruning. In 1.0 a refutation is a verdict that scrolls past; in 2.0 it closes a branch, promotes an alternative under a stated rule, and stays in a record that the next cycle consults. Pruning becomes cumulative instead of repeated.

It makes the record auditable by someone who was not there. Every claim carries a validator, its hash, the model that wrote it, the reviewer’s objection, the oracle’s output, and the decision that followed. That is the difference between a transcript and evidence, and it is what makes an LLM-driven result quotable at all.

It makes the system’s own failures visible, and separable. This is the deepest argument and the one the comparison benchmark cannot score, so it is made from live use. A campaign *acts* on its own signals, and doing so exposes a class of defect 1.0 cannot exhibit: *infrastructure mistaken for a research signal*. Six were found and fixed while a single research programme ran across two campaigns and more than a dozen episodes.

Transport failure	Was being scored as
a provider 500 mid-pipeline	an uninformative research episode (it paused a branch early)
a branch blind to its ancestor’s result	a direction that had to re-derive what was already certified
a malformed repair patch	a reason to destroy an approved, passing validator
a reviewer objection lost between episodes	a fresh author reinventing the rejected shortcut
a Sage <code>sys.exit(0)</code> on success	an operational error masking a printed PASS
a cold WSL boot exceeding a 10s probe	an installed engine reported as missing

Each was fixed at its source. The point is not the list but its signature: every one is the five axes doing their job—keeping a failure of *transport* from being recorded as a failure of *science*. 1.0 has no place for these to happen, and equally no place to catch them; the layer that creates the risk is the layer that makes it visible.

6.1 A worked example: a research programme, end to end

The clearest evidence that the campaign layer buys something is that it drove a real physics question to certified verdicts of *both* signs, auditable afterwards. The question—whether a modification to a warp ansatz repairs a pointwise energy-condition violation—produced, in sequence:

- a **certified positive**: the modification makes the condition hold over an entire null cone at one point, exact rational certificate, two independent methods agreeing;
- a **certified negative**: at another point the same modification *fails*, with an exact witness direction—the fix is local, not general.

Two features of that arc are exactly what 2.0 exists to provide. First, the negative result is a first-class deliverable, not a discarded branch: it closes the question (one counterexample refutes the universal) and it is recorded with its witness. Second, one episode showed the axes preventing a false record in real time: an ensemble stated a constant wrong by one part in 10^{11} , a validator refuted *that* constant, and the re-anchoring axis marked the verdict SUBSTITUTED—so a transcription slip was recorded as a refuted *substituted* claim rather than as a refuted *scientific* one. In 1.0, with a single verdict field, that same episode reads as a refutation of the physics. That difference—between a wrong headline and a correct one—is the whole thesis of the second version, caught happening.

7 When to use each line

The two are not rivals; 2.0 wraps 1.0, so the choice is whether a given task needs the wrapper.

Reach for 1.0 (the bare atomic cycle) when	Reach for 2.0 (the campaign layer) when
the question is a single decidable proposition	the question is a programme—many linked claims, branches, dead ends
you will read the verdict yourself, now	someone who was not there must audit the result later
latency matters more than a durable record	a false acceptance would be expensive to discover downstream
you are smoke-testing a validator or an engine	negative results must accumulate rather than be re-derived
throughput per unit of quota is the priority	the cost of being confidently wrong exceeds the cost of a timeout

Because 2.0’s per-case latency is $1.75 \times$ 1.0’s and its accuracy advantage is not yet significant, the honest default for a one-off decidable check is 1.0. The case for 2.0 is not speed or raw accuracy; it is that the result is *safe to quote*—no false acceptances, a full audit trail, and a verdict whose axes will not let an operational hiccup masquerade as a scientific finding.

8 Position

ASTRA 2.0 is a strict superset of 1.0 in behaviour, with 9 new modules, three times the test coverage, nine epistemic controls each traceable to a measured failure, and a persistent evidence ledger with guarantees 1.0 has no need for. It is demonstrably more *careful*.

Whether it is more *productive* per unit of quota is not established, and this document declines to assert it until the paired benchmark lands and a blinded expert scoring of trajectories exists. The architecture report states the same thing about its internal ablation. Claiming otherwise would be exactly the failure mode the whole system is built to prevent.

Note on scope. Both lines run on commercially available subscription models from three providers on one workstation (Section 2). Nothing here is a claim about frontier or unreleased systems, and no result depends on privileged model access.